

Hardware Transactional Memory (2)

Mohamed Mohamedin

Chapter 5 of TM Book

HTM Limitation

- Best-efforts
 - Transactions are not guaranteed to commit
 - Must provide a software fallback path

Non-Speculative Fallback

```
if (_xbegin() == _XBEGIN_STARTED) {  
    read lock state  
    if (lock taken) _xabort();  
    work;  
    _xend()  
} else {  
    lock->lock();  
    work;  
    lock->unlock();  
}
```

Non-Speculative Fallback

```
if ( _xbegin() == _XBEGIN_STARTED) {  
    read lock state  
    if (lock taken) _xabort();  
    work;  
    _xend()  
}
```

**reading lock ensures that
transaction will abort if another
thread acquires lock**

Non-Speculative Fallback

```
if (_xbegin() == _XBEGIN_STARTED) {  
    read lock state  
    if (lock taken) _xabort();  
    work;  
    _xend()  
} else {  
    lock state  
    work  
    lock state  
}
```

**abort if another thread has
acquired lock**


Non-Speculative Followup

**on abort, acquire lock & do work
(aborting concurrent speculative
transactions)**

```
if  
re  
if (lock taken) _xabort();  
work;  
_xend()  
} else {  
  lock->lock();  
  work;  
  lock->unlock();  
}
```

Global Lock Fallback


- Must define a software fallback path
 - Default is global lock



- `_xbegin()`
- `if (read(lock)) == 1 then _xabort()`

Global Lock Fallback


- Must define a software fallback path
 - Default is global lock



- `_xbegin()`
- `if (read(lock)) == 1 then _xabort()`
- do some work

Global Lock Fallback

- Must define a software fallback path
 - Default is global lock



- `_xbegin()`
- `if (read(lock)) == 1 then _xabort()`
- do some work

`//non-transactionally
CAS(lock, 0, 1)`

Global Clock Fallback


- Pros
 - Very easy to program
 - 100% safe
- Cons
 - Unnecessary aborts concurrent HTM transactions
 - Limit concurrency
 - Lemming effect
 - Cascading the failures in HTM, ending up with all transactions running in the global lock fallback path

Enhancing the Fallback Path

- Target
 - Enhance concurrency
 - Allow non-conflicting HTM transaction to run with the fallback path
 - Low overhead on HTM
 - An instrumentation is needed to communicate between HTM and the fallback path
 - E.g., The Global Lock
 - A heavy instrumentation would annul the advantages of HTM
 - Allow most transactions to commit in HW
 - No lemming effect

Lemming Effect

- Once a transaction take the fallback path → All remaining transactions follows
- How?



Lemming Effect

- One a transaction acquire the lock → all concurrent transactions abort

```
if (_xbegin() == _XBEGIN_STARTED) {  
    read lock state  
    if (lock taken) _xabort();  
    work;  
    _xend()  
} else {  
    lock->lock();  
    work;  
    lock->unlock();  
}
```

Lock is in the read-set of the HTM transaction, then when it is locked in the fallback path (a write), a conflict aborts the HTM transaction

Lemming Effect

- Upon immediate retry, all HTM transactions will abort since the lock is already acquired

```
if (_xbegin() == _XBEGIN_STARTED) {  
    read lock state  
    if (lock taken) _xabort();  
    work;  
    _xend()  
} else {  
    lock->lock();  
    work;  
    lock->unlock();  
}
```

Lock is already taken, so all HTM retries will abort immediately

Lemming Effect

- Hence, all HTM transaction will take the fallback path

```
if (_xbegin() == _XBEGIN_STARTED) {  
    read lock state  
    if (lock taken) _xabort();  
    work;  
    _xend()  
} else {  
    lock->lock();  
    work;  
    lock->unlock();  
}
```

All HTM transaction will fallback to the global lock

Lemming Effect

- This behavior will continue until the next quiescent time

```
if (_xbegin() == _XBEGIN_STARTED) {  
    read lock state  
    if (lock taken) _xabort();  
    work;  
    _xend()  
} else {  
    lock->lock();  
    work;  
    lock->unlock();  
}
```


Lemming Effect

- Eliminating lemming effect is easy
 - Wait for the lock to be free before retrying

```
while (lock taken) wait();  
if (_xbegin() == XBEGIN_STARTED) {  
    read lock state;  
    if (lock taken) _xabort();  
    work;  
    _xend()  
} else {  
    lock->lock();  
    work;  
    lock->unlock();  
}
```

Wait until lock is free before trying an HTM transaction again

Lazy Subscription

- Move global lock checking to the end of a transaction

```
if (_xbegin() == _XBEGIN_STARTED) {  
    //Do transaction work first  
    work;  
    //Check lock just before committing  
    if (lock taken) _xabort();  
    _xend()  
} else {  
    lock->lock();  
    work;  
    lock->unlock();  
}
```

Lazy Subscription

- Why this is beneficial?
 - It allows more concurrency

```
if (_xbegin() == _XBEGIN_STARTED) {  
    //Do transaction work first  
    work;  
    //Check lock just before committing  
    if (lock taken) _xabort();  
    _xend()  
} else {  
    lock->lock();  
    work;  
    lock->unlock();  
}
```

Lazy Subscription

- Concurrent transactions will not immediately abort

```
if (_xbegin() == _XBEGIN_STARTED) {  
    //Do transaction work first  
    work;  
    //Check lock just before committing  
    if (lock taken) _xabort();  
    _xend()  
} else {  
    lock->lock();  
    work;  
    lock->unlock();  
}
```

Lazy Subscription

- New transactions will be allowed to start while the global lock is acquired

```
if (_xbegin() == _XBEGIN_STARTED) {  
    //Do transaction work first  
    work;  
    //Check lock just before committing  
    if (lock taken) _xabort();  
    _xend()  
} else {  
    lock->lock();  
    work;  
    lock->unlock();  
}
```

Lazy Subscription

- What about safety?
 - HTM transactions can access intermediate changes!

```
if (_xbegin() == _XBEGIN_STARTED) {  
    //Do transaction work first  
    work;  
    //Check lock just before committing  
    if (lock taken) _xabort();  
    _xend()  
} else {  
    lock->lock();  
    work;  
    lock->unlock();  
}
```

Lazy Subscription

- What about safety?
 - HTM is sandboxed and exceptions or infinite loops will the transaction

```
if (_xbegin() == _XBEGIN_STARTED) {  
    //Do transaction work first  
    work;  
    //Check lock just before committing  
    if (lock taken) _xabort();  
    _xend()  
} else {  
    lock->lock();  
    work;  
    lock->unlock();  
}
```

Lazy Subscription

- What about safety?
 - Unfortunately, current HTM implementation is not 100% sandboxed

```
if (_xbegin() == _XBEGIN_STARTED) {  
    //Do transaction work first  
    work;  
    //Check lock just before committing  
    if (lock taken) _xabort();  
    _xend()  
} else {  
    lock->lock();  
    work;  
    lock->unlock();  
}
```


Lazy Subscription Discussion




Figure 1. Obvious SGL Fallback implementation (E-SGL).




Figure 2. Lazy SGL (L-SGL).

Lazy Subscription Discussion

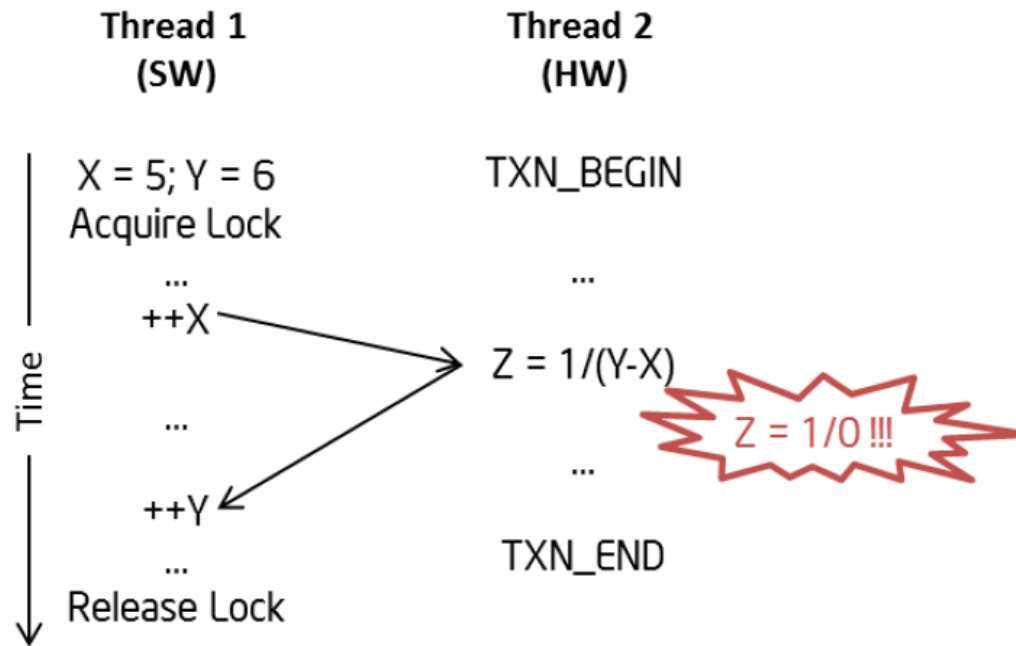


Figure 3. Inconsistent reads.

Lazy Subscription Discussion

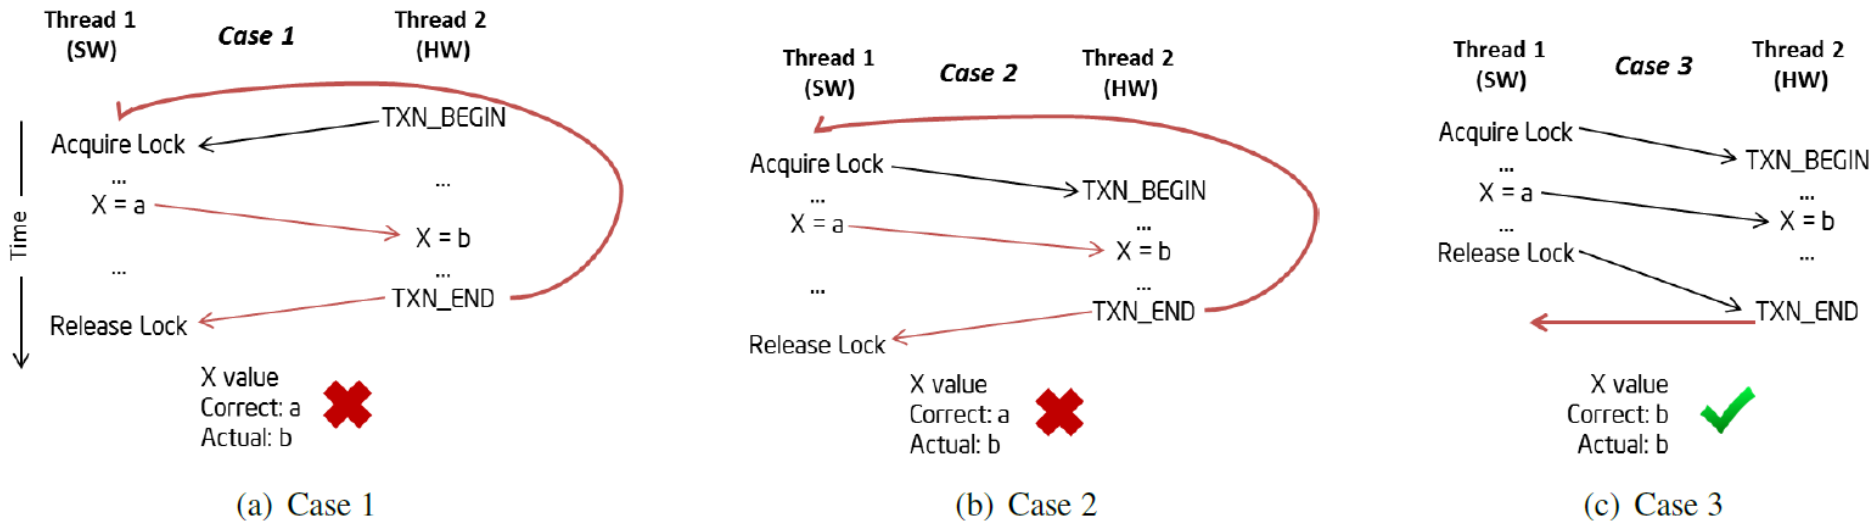


Figure 4. Correctness: Cases 1-3. Arrows denote the “happens-before” relation.

Lazy Subscription Discussion

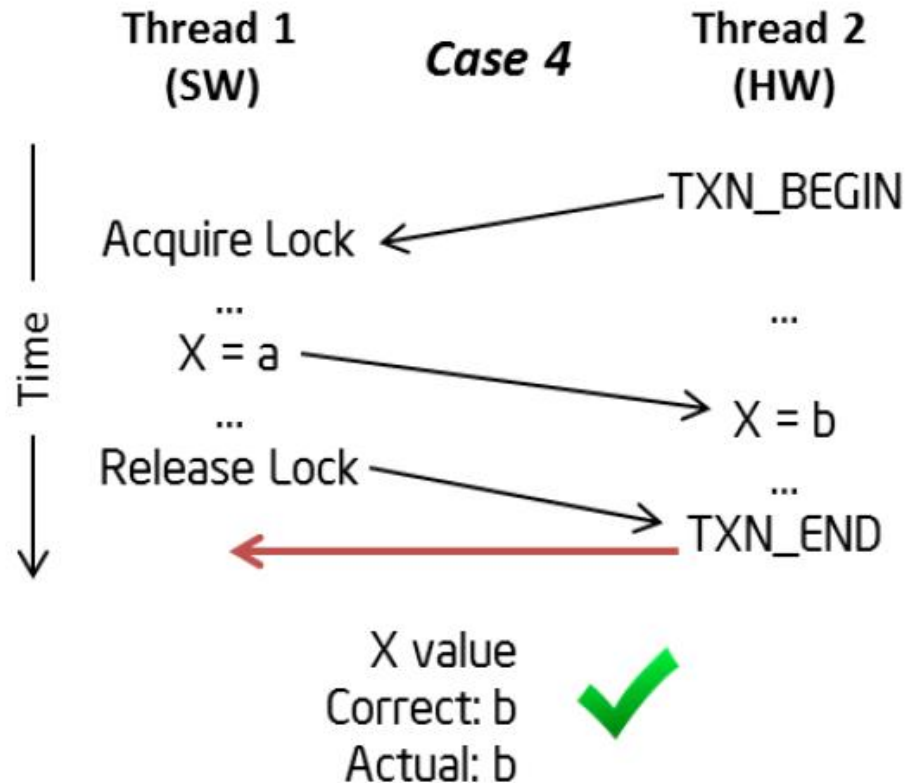


Figure 5. Correctness: Case 4.

Pitfalls of Lazy Subscription

- Observing inconsistent state

```
1 void (*method_table[2])() = {method1, method2};

3 int next_method = 0;

5 lock L;

7 void apply_next() {
8     acquire(L);
9     (*method_table[next_method])();
10    if (++next_method > 2)
11        next_method = 0;
12    release(L);
13 }
```

Pitfalls of Lazy Subscription

- Indirect branch

```
1 void (*method_table[2])() = {method1, method2};

3 int next_method = 0;

5 lock L;

7 void apply_next() {
8     acquire(L);
9     (*method_table[next_method])();
10    if (++next_method > 2)
11        next_method = 0;
12    release(L);
13 }
```

Pitfalls of Lazy Subscription

- Lock scribbling
- Subscribing to the wrong “lock”
- Self modifying code
- Corrupted return address

Pitfalls of Lazy Subscription

- Conclusion
 - Lazy Subscription is too risky on current architectures
 - HW must be modified first to support complete sandboxing

Retrying HTM Transactions

- When should we retry?
- How many retrials are enough?
- Should we try HTM in the first place?

Retrying HTM Transactions

Algorithm 1 TSX in GCC

```
1: int attempts  $\leftarrow$  2
2: int status  $\leftarrow$  XBEGIN
3: if status  $\neq$  ok then
4:   if attempts = 0 then
5:     acquire(globalLock)
6:   else
7:     attempts--
8:     goto line 2
9:   if is_locked(globalLock)
10:    XABORT
11:  $\triangleright$  ...transactional code
12: if attempts = 0 then
13:   release(globalLock)
14: else
15:   XEND
```

Abort code
retry : Transient failure
conflict : Contention to data
capacity : Exceeded cache
explicit : XABORT invoked
other : Faults, preemption, ...

Fig. 2. Error codes in TSX.

*Self-Tuning Intel Transactional
Synchronization Extensions

Retrying HTM Transactions

Algorithm 1 TSX in GCC

```
1: int attempts  $\leftarrow$  2
2: int status  $\leftarrow$  XBEGIN
3: if status  $\neq$  ok then
4:   if attempts = 0 then
5:     acquire(globalLock)
6:   else
7:     attempts--
8:     goto line 2
9:   if is_locked(globalLock)
10:    XABORT
11:  $\triangleright$  ...transactional code
12: if attempts = 0 then
13:   release(globalLock)
14: else
15:   XEND
```

This code ignores the status code and just retry without considering the reason of the abort

capacity:	Exceeded cache
explicit:	XABORT invoked
other:	Faults, preemption, ...

Fig. 2. Error codes in TSX.

*Self-Tuning Intel Transactional Synchronization Extensions

Retrying HTM Transactions

Algorithm 2 HEURISTIC based approach for TSX.

```
1: int attempts  $\leftarrow$  5
    $\triangleright$  avoid the lemming effect
2: while(is_locked(global-lock)) do pause  $\triangleright$  x86 instruction
3: int status  $\leftarrow$  XBEGIN
4: if status  $\neq$  ok then
5:   if attempts = 0 then
6:     acquire(global-lock)
7:   else
8:     if status = explicit  $\vee$  status = other then
9:       attempts  $\leftarrow$  attempts - 1  $\triangleright$  skipped if transient
10:    else if status = capacity then
11:      attempts  $\leftarrow$  0  $\triangleright$  give up, likely that it always fails
12:      goto line 2
13:  $\triangleright$  ...code to run in transaction
14: if attempts = 0 then
15:   release(global-lock)
16: else
17:   if is_locked(global-lock) then
18:     XABORT  $\triangleright$  check for concurrent pessimistic thread
19:   XEND
```

*Self-Tuning Intel Transactional Synchronization Extensions

Retrying HTM Transactions

Algorithm 2 HEURISTIC based approach for TSX.

```
1: int attempts ← 5
   ▷ avoid the lemming effect
2: while(is_locked(global-lock)) do pause                                ▷ x86 instruction
3: int status ← XBEGIN
4: if status ≠ ok then
5:   if attempts = 0 then
6:     acquire(global-lock)
7:   else
8:     if status = explicit ∨ status = capacity then
9:       attempts ← attempts - 1
10:    else if status = capacity then
11:      attempts ← 0
12:      goto line 2
13: ▷ ...code to run in transaction
14: if attempts = 0 then
15:   release(global-lock)
16: else
17:   if is_locked(global-lock) then
18:     XABORT                                                                ▷ check for concurrent pessimistic thread
19:   XEND
```

Avoid lemming effect by waiting for global lock to be free before attempting an HTM transaction

▷ skipped if transient

▷ give up, likely that it always fails

*Self-Tuning Intel Transactional Synchronization Extensions

Retrying HTM Transactions

Algorithm 2 HEURISTIC based approach for TSX.

```
1: int attempts ← 5
   ▷ avoid the lemming effect
2: while(is_locked(global-lock)) do pause
3: int status ← XBEGIN
4: if status ≠ ok then
5:   if attempts = 0 then
6:     acquire(global-lock)
7:   else
8:     if status = explicit ∨ status = other then
9:       attempts ← attempts - 1
10:    else if status = capacity then
11:      attempts ← 0
12:      goto line 2
13: ▷ ...code to run in transaction
14: if attempts = 0 then
15:   release(global-lock)
16: else
17:   if is_locked(global-lock) then
18:     XABORT
19:   XEND
```

If the abort is explicit or for “other” reasons, consume one attempt

▷ skipped if transient

▷ give up, likely that it always fails

▷ check for concurrent pessimistic thread

*Self-Tuning Intel Transactional Synchronization Extensions

Retrying HTM Transactions

Algorithm 2 HEURISTIC based approach for TSX.

```
1: int attempts  $\leftarrow$  5
    $\triangleright$  avoid the lemming effect
2: while(is_locked(global-lock)) do pause  $\triangleright$  x86 instruction
3: int status  $\leftarrow$  XBEGIN
4: if status  $\neq$  ok then
5:   if attempts = 0 then
6:     acquire(global-lock)
7:   else
8:     if status = explicit  $\vee$  status = other then
9:       attempts  $\leftarrow$  attempts - 1  $\triangleright$  skipped if transient
10:    else if status = capacity then
11:      attempts  $\leftarrow$  0
12:      goto line 2
13:  $\triangleright$  ...code to run in transaction
14: if attempts = 0 then
15:   release(global-lock)
16: else
17:   if is_locked(global-lock) then
18:     XABORT  $\triangleright$  check for concurrent pessimistic thread
19:   XEND
```

Capacity aborts consume all attempts as there is no benefit in retrying a transaction that doesn't fit in HTM that it always fails

*Self-Tuning Intel Transactional Synchronization Extensions

Retrying HTM Transactions

Algorithm 2 HEURISTIC based approach for TSX.

```
1: int attempts  $\leftarrow$  5
    $\triangleright$  avoid the lemming effect
2: while(is_locked(global-lock)) do pause
3: int status  $\leftarrow$  XBEGIN
4: if status  $\neq$  ok then
5:   if attempts = 0 then
6:     acquire(global-lock)
7:   else
8:     if status = explicit  $\vee$  status = other then
9:       attempts  $\leftarrow$  attempts - 1
10:    else if status = capacity then
11:      attempts  $\leftarrow$  0
12:    goto line 2
13:  $\triangleright$  ...code to run in transaction
14: if attempts = 0 then
15:   release(global-lock)
16: else
17:   if is_locked(global-lock) then
18:     XABORT
19:   XEND
```

\triangleright x86 instruction

Conflict & Retry aborts do NOT consume any attempts as these are normal expected TM aborts

\triangleright check for concurrent pessimistic thread

*Self-Tuning Intel Transactional Synchronization Extensions

Retrying HTM Transactions

Algorithm 2 HEURISTIC based approach for TSX.

```
1: int attempts  $\leftarrow$  5
    $\triangleright$  avoid the lemming effect
2: while(is_locked(global-lock)) do pause  $\triangleright$  x86 instruction
3: int status  $\leftarrow$  XBEGIN
4: if status  $\neq$  ok then
5:   if attempts = 0 then
6:     acquire(global-lock)
7:   else
8:     if status = explicit  $\vee$  status = capacity then
9:       attempts  $\leftarrow$  attempts - 1  $\triangleright$  skipped if transient
10:    else if status = capacity then
11:      attempts  $\leftarrow$  0  $\triangleright$  give up, likely that it always fails
12:      goto line 2
13:  $\triangleright$  ...code to run in transaction
14: if attempts = 0 then
15:   release(global-lock)
16: else
17:   if is_locked(global-lock) then  $\triangleright$  check for concurrent pessimistic thread
18:     XABORT
19:   XEND
```

When all attempts are consumed,
acquire the global lock

*Self-Tuning Intel Transactional Synchronization Extensions

Retrying HTM Transactions

Algorithm 2 HEURISTIC based approach for TSX.

```
1: int attempts  $\leftarrow$  5
    $\triangleright$  avoid the lemming effect
2: while(is_locked(global-lock)) do pause  $\triangleright$  x86 instruction
3: int status  $\leftarrow$  XBEGIN
4: if status  $\neq$  ok then
5:   if attempts = 0 then
6:     acquire(global-lock)
7:   else
8:     if status = explicit  $\vee$  status = other then
9:       attempts  $\leftarrow$  attempts - 1  $\triangleright$  skipped if transient
10:    else if status = capacity then
11:      attempts  $\leftarrow$  0  $\triangleright$  give up, likely that it always fails
12:      goto line 2
13:  $\triangleright$  ...code to run in transaction
14: if attempts = 0 then
15:   release(global-lock)
16: else
17:   if is_locked(global-lock) then  $\triangleright$  for concurrent pessimistic thread
18:     XABORT
19:   XEND
```

If attempt are zero, then I must have acquired the lock and need to unlock it

*Self-Tuning Intel Transactional Synchronization Extensions

Retrying HTM Transactions

Algorithm 2 HEURISTIC based approach for TSX.

```
1: int attempts  $\leftarrow$  5
    $\triangleright$  avoid the lemming effect
2: while(is_locked(global-lock)) do pause  $\triangleright$  x86 instruction
3: int status  $\leftarrow$  XBEGIN
4: if status  $\neq$  ok then
5:   if attempts = 0 then
6:     acquire(global-lock)
7:   else
8:     if status = explicit  $\vee$  status = other then
9:       attempts  $\leftarrow$  attempts - 1  $\triangleright$  skipped if transient
10:    else if status = capacity then
11:      attempts  $\leftarrow$  0  $\triangleright$  p, likely that it always fails
12:      goto line 2
13:  $\triangleright$  ...code to run in transaction
14: if attempts = 0 then
15:   release(global-lock)
16: else
17:   if is_locked(global-lock) then
18:     XABORT  $\triangleright$  check for concurrent pessimistic thread
19:   XEND
```

Lazy subscription to the lock and
_xend if lock is free

*Self-Tuning Intel Transactional Synchronization Extensions

Retrying HTM Transactions

Algorithm 3 TUNER adaptive configuration.

```
1: int ucbBelief  $\leftarrow$   $\triangleright$  last configuration used
2: int attempts  $\leftarrow$   $\triangleright$  last configuration used
3: if reoptimize() then
4:   long initCycles  $\leftarrow$  obtainRDTSC()
5:   while is_locked(global-lock) do pause
6:   int status  $\leftarrow$  XBEGIN
7:   if status  $\neq$  ok then
8:     if attempts = 0 then
9:       if reoptimize() then tuneAttempts(ucbBelief)
10:      acquire(global-lock)
11:    else
12:      if status = capacity then
13:         $\triangleright$  set attempts according to ucbBelief
14:      else if status = explicit  $\vee$  status = other then
15:        attempts  $\leftarrow$  attempts - 1
16:        goto line 5
17:  $\triangleright$  ...code to run in transaction
18: if attempts = 0 then
19:   release(global-lock)
20: else
21:   if is_locked(global-lock) then XABORT
22:   XEND
23: if reoptimize() then
24:   long totalCycles  $\leftarrow$  obtainRDTSC() - initCycles
25:   ucbBelief  $\leftarrow$  UCB(totalCycles)
26:   attempts  $\leftarrow$  GRAD(totalCycles)
```

*Self-Tuning Intel Transactional
Synchronization Extensions

\triangleright rules of Section 4.2.
 \triangleright rules of Section 4.3.

Retrying HTM Transactions

Algorithm 3 TUNER adaptive configuration.

```
1: int ucbBelief  $\leftarrow$   $\triangleright$  last configuration used
2: int attempts  $\leftarrow$   $\triangleright$  last configuration used
3: if reoptimize() then
4:   long initCycles  $\leftarrow$  obtainRDTSC()
5:   while is_locked(global-lock) do pause
6:   int status  $\leftarrow$  XBEGIN
7:   if status  $\neq$  ok then
8:     if attempts = 0 then
9:       if reoptimize() then tuneAttempts(ucbBelief)
10:      acquire(global-lock)
11:    else
12:      if status = capacity then
13:         $\triangleright$  set attempts according to ucbBelief
14:      else if status = explicit  $\vee$  status = other then
15:        attempts  $\leftarrow$  attempts - 1
16:        goto line 5
17:   $\triangleright$  ...code to run in transaction
18:  if attempts = 0 then
19:    release(global-lock)
20:  else
21:    if is_locked(global-lock) then XABORT
22:    XEND
23:  if reoptimize() then
24:    long totalCycles  $\leftarrow$  obtainRDTSC() - initCycles
25:    ucbBelief  $\leftarrow$  UCB(totalCycles)
26:    attempts  $\leftarrow$  GRAD(totalCycles)
```

Number of attempts is adaptive

*Self-Tuning Intel Transactional
Synchronization Extensions

\triangleright rules of Section 4.2.
 \triangleright rules of Section 4.3.

Retrying HTM Transactions

Algorithm 3 TUNER adaptive configuration.

```
1: int ucbBelief  $\leftarrow$   $\triangleright$  last configuration used
2: int attempts  $\leftarrow$   $\triangleright$  last configuration used
3: if reoptimize() then
4:   long initCycles  $\leftarrow$  obtainRDTSC()
5:   while is_locked(global-lock) do pause
6:   int status  $\leftarrow$  XBEGIN
7:   if status  $\neq$  ok then
8:     if attempts = 0 then
9:       if reoptimize() then tuneAttempts(ucbBelief)
10:      acquire(global-lock)
11:    else
12:      if status = capacity then
13:         $\triangleright$  set attempts according to ucbBelief
14:      else if status = explicit  $\vee$  status = other then
15:        attempts  $\leftarrow$  attempts - 1
16:      goto line 5
17:  $\triangleright$  ...code to run in transaction
18: if attempts = 0 then
19:   release(global-lock)
20: else
21:   if is_locked(global-lock) then XABORT
22:   XEND
23: if reoptimize() then
24:   long totalCycles  $\leftarrow$  obtainRDTSC() - initCycles
25:   ucbBelief  $\leftarrow$  UCB(totalCycles)
26:   attempts  $\leftarrow$  GRAD(totalCycles)
```

Consuming all attempts on capacity aborts is not always the best option

*Self-Tuning Intel Transactional Synchronization Extensions

\triangleright rules of Section 4.2.
 \triangleright rules of Section 4.3.

Global Lock Fallback

- Even with all these optimizations
 - Once the global clock is acquired, no concurrency is allowed
 - A long running transaction acquiring the global clock will block the progress of all other transactions
 - Without having any mutual conflicts
- Can we do better?

Hybrid TM

- Mix HTM with STM
- Allows both HTM and STM to work concurrently (if possible)
 - Some proposals (e.g., PhaseTM) work in phases where all transactions run in HTM or STM
- Requires some instrumentation inside HTM path

Hybrid TM

- One naïve solution is to use per location metadata (Orec based STM)
 - For each read in HTM, check the associated metadata
 - Writes must acquire the associated locks to inform STM transactions of that write

Hybrid TM

- One naïve solution is to use per location metadata (Orec based STM)
 - The resulting algorithm performance is very bad
 - Consumes a lot of the limited HTM resources
 - Extra reads and writes due to metadata manipulation
 - Most HTM transactions abort
 - False aborts due to space limitation (metadata access)
 - False conflict due to global shared metadata access (e.g., global clock)
- What's the solution?

Hybrid TM

- Minimal instrumentation
- Minimal communication between HTM and STM
- For example, do NOT use Orecs

Hybrid TM

- Minimal instrumentation
 - Minimal communication between HTM and STM
 - For example, do NOT use Orecs
 - Best candidate from algorithms we learned
- NRec

Hybrid NOrec

- NOrec
 - A single shared global lock is used
 - Minimal communication between STM and HTM
 - Value-based validation
 - No per-location metadata
 - Minimal instrumentation inside HTM

Hybrid NOrec

- NOrec STM Algorithm

```
1  padded unsigned seqlock
3  thread local unsigned snapshot
4  thread local ReadSet reads
5  thread local WriteSet writes

6  SW_VALIDATE
7    snapshot = seqlock
8    if (snapshot & 1)
9      goto 7
10   foreach (addr, val) in reads
11     if (*addr != val)
12       SW_ABORT
13   if (snapshot != seqlock)
14     goto 7

15  SW_BEGIN
16    snapshot = seqlock
17    if (snapshot & 1)
18      goto 16

19  SW_COMMIT
20    if (writes.empty())
21      return
22    while (!CAS(&seqlock, snapshot,
23              snapshot + 1))
24      SW_VALIDATE
25    foreach (addr, val) in writes
26      *addr = val
27    seqlock = seqlock + 1
28    reads.reset(), writes.reset()

28  SW_READ(addr)
29    if (addr in writes)
30      return writes.find(addr)
31    val = *addr
32    if (snapshot != seqlock)
33      SW_VALIDATE
34      goto 31
35    reads.append(addr, val)
36    return val

38  SW_WRITE(addr, val)
39    writes.append(addr, val)

41  SW_ABORT
42    reads.reset(), writes.reset()
43    /* restart transaction */
```

*Hybrid NOrec: A Case Study in the Effectiveness of Best Effort Hardware Transactional Memory

Hybrid NOrec

- Simply, subscribe to NOrec global lock at the beginning
- Increment NOrec global lock at the end to notify STM transaction of the change

```
5    HW_POST_BEGIN
6        if (seqlock & 1)
7            while (true) // await abort

9    HW_PRE_COMMIT
10       seqlock = seqlock + 2
```

Hybrid NOrec

- Simply, subscribe to NOrec global lock at the beginning
- Increment NOrec global lock at the end to notify STM transaction of the change

```
5  HW_POST_BEGIN
6      if (seqlock & 1)
7          while (true)
8              // ...
9  HW_PRE_COMMIT
10     seqlock = seqlock + 2
```

All HTM reading seqlock will abort
one this line is executed!

Hybrid NOrec

- This improvement, reduces HTM-HTM false conflicts due to writing to seqlock

- Shorter window of conflicts

```
1  padded unsigned seqlock
2  padded unsigned counter
```

```
5  HW_POST_BEGIN
6    if (seqlock & 1)
7      while (true) // await abort
```

```
9  HW_PRE_COMMIT
10 counter = counter + 1
```

Hybrid NOrec

- This improvement, reduces HTM-HTM false conflicts due to writing to seqlock

- Shorter window of conflicts

```
1  padded unsigned seqlock
2  padded unsigned counter
```

```
5  HW_POST_BEGIN
6    if (seqlock & 1)
7      while (true) //
9  HW_PRE_COMMIT
10   counter = counter + 1
```

With using two locations for communication. Only HTM transactions executing this line concurrently will conflict with each other

Hybrid NOrec

- This improvement, eliminates HTM-HTM false conflicts due to writing to same counter

```
1  padded unsigned seqlock
2  padded unsigned counter[]
3  thread local unsigned id

5  HW_POST_BEGIN
6      if (seqlock & 1)
7          while (true) // await abort

9  HW_PRE_COMMIT
10     counter[id] = counter[id] + 1
```

Hybrid NOrec

- This improvement, eliminates HTM-HTM false conflicts due to writing to same counter

```
1  padded unsigned seqlock
2  padded unsigned counter[]
3  thread local unsigned id

5  HW_POST_BEGIN
6      if (seqlock & 1)
7          while (true) /
9  HW_PRE_COMMIT
10     counter[id] = counter[id] + 1
```

Now, each thread has its own location to indicate its writes to STM transactions, so no more false conflicts due to this communication

Hybrid NOrec

- This improvement, eliminates HTM-HTM false conflicts due to writing to same counter

```
1  padded unsigned seqlock
2  padded unsigned counter[]
3  thread local unsigned id

5  HW_POST_BEGIN
6      if (seqlock & 1)
7          while (true) /
9  HW_PRE_COMMIT
10     counter[id] = counter[id] + 1
```

BUT, software tx_commit will have more overhead checking all these counters